

JADAVPUR UNIVERSITY

**Department
Of
Computer Science and Engineering**

NAME : SUBHADIP DUTTA

MCA 2nd YEAR 2nd SEMESTER

ROLL NO : 002410503054

REG NO : 1665857 OF 2024-25

TOPIC : TRANSFORMERS IN NLP

ABSTRACT

Natural Language Processing (NLP) is a rapidly evolving field of Artificial Intelligence that focuses on enabling computers to understand, interpret, and generate human language. The primary goal of NLP is to bridge the communication gap between humans and machines by allowing systems to process textual and spoken language efficiently. Applications of NLP include machine translation, text summarization, sentiment analysis, question answering, speech recognition, and conversational AI systems such as chatbots and virtual assistants.

Early NLP systems relied heavily on traditional machine learning techniques and sequential neural architectures like **Recurrent Neural Networks (RNNs)**. Although RNNs were capable of processing sequential data, they suffered from several limitations, including vanishing and exploding gradient problems, difficulty in capturing long-range dependencies, and slow sequential computation. **Long Short-Term Memory (LSTM)** networks were introduced to partially address these issues through memory cells and gating mechanisms, enabling better retention of contextual information over longer sequences. However, LSTMs still faced challenges related to parallelization, computational efficiency, and maintaining dependencies in very long sentences.

The introduction of the Transformer architecture revolutionized NLP by eliminating recurrence entirely and replacing it with the self-attention mechanism. Unlike RNNs and LSTMs, Transformers process all tokens in a sequence simultaneously, enabling highly parallel computation and significantly improving training efficiency. This seminar explores the Transformer architecture in detail, beginning with the motivation behind its development and the shortcomings of previous sequential models.

The seminar provides an in-depth explanation of the core components of **Transformers**, including self-attention, multi-head attention, positional encoding, feed-forward neural networks, encoder-decoder architecture, and masked attention mechanisms. Special emphasis is placed on understanding why self-attention is needed, how it captures global contextual relationships between words, and how it overcomes the limitations of RNNs and LSTMs. The role of positional encoding in preserving word order information is also discussed in detail.

Furthermore, here explains the concept of multi-head attention and masked multi-head attention, highlighting their differences and importance in tasks such as language modeling and text generation. The complete training and testing process of Transformers is also examined to provide a practical understanding of how these models learn and generate meaningful language representations.

In addition, here discusses the major advantages of Transformers, such as parallel processing, improved contextual understanding, scalability, and superior performance on large datasets, along with their limitations, including high computational requirements and large memory consumption. Finally, we explore several popular Transformer-based models, including **BERT**, **GPT**, **CLAUDE** and **LLaMA**, which have significantly advanced the state of the art in NLP applications.

This seminar aims to provide a comprehensive and accessible understanding of Transformer architectures in NLP, making it suitable for students, researchers, and practitioners interested in modern language processing systems and the future of Artificial Intelligence.